

04a_exploratory_pandas

June 24, 2015

1 Getting Started with Exploratory Data Analysis

3 important Python packages 1. NumPy for efficient computation on arrays 2. Pandas for data analysis 3. Matplotlib for plotting in the notebook

```
In [2]: import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
```

1.1 Pandas

Python module for manipulating tabular data

1.2 pandas

- Provides python a **DataFrame**
- Structured manipulation tools
- Built on top of **numpy**
- Huge growth from 2011-2012
- Very **efficient**
- Great for *medium* data

Resources

- pandas.pydata.org
- [Python for Data Analysis](#) by Wes McKinney
- [Data Wrangling Kung Fu with Pandas](#) by Wes McKinney
- [Cheat sheet](#) by Quandl

1.2.1 Why pandas?

80% of the effort in data analysis is spent cleaning data. [Hadley Wickham](#)

Efficiency

- Different views of data
- [Tidy data](#) by Hadley Wickham

Raw data is often in the wrong format

- How often to you download an array ready for array-oriented computing?
- e.g. `scikit-learn` interface

Storage may be best in a different format

- Sparse representations
- Upload to database

1.3 Simple example using the Walrus data

1.3.1 Reading a CSV file as text

```
In [3]: filename = os.path.join('Walrus_Data', 'Walruses.csv')
        f = open(filename, 'r')
        lines = f.readlines()
        lines[:10]
```

```
Out[3]: ['Walrus,DateTimeUTC,Xcoord,Ycoord,Behav,Longitude,Latitude\r\n',
         '271,5/31/2008 19:25,"95,616.95",-528,324.60",1.009,-167.9560949,65.24871506\r\n',
         '271,6/1/2008 3:24,"84,741.71",-511,653.75",1.0005,-168.1779869,65.40121711\r\n',
         '271,6/1/2008 11:24,"71,834.45",-491,176.95",1.00625,-168.4443605,65.58796906\r\n',
         '271,6/1/2008 19:24,"65,275.80",-478,935.62",1.02025,-168.5802843,65.6991429\r\n',
         '271,6/2/2008 3:24,"69,343.24",-473,948.91",1.00775,-168.4892154,65.74298362\r\n',
         '271,6/2/2008 11:24,"72,634.53",-457,308.67",1,-168.4082441,65.89142326\r\n',
         '271,6/2/2008 19:24,"73,253.86",-425,586.14",1,-168.3764825,66.17566325\r\n',
         '271,6/3/2008 3:24,"79,223.97",-401,784.87",1.00975,-168.2291622,66.38754225\r\n',
         '271,6/3/2008 11:24,"77,052.23",-382,920.49",1.20225,-168.2658694,66.5571914\r\n']
```

1.3.2 Creating a DataFrame

```
df = pd.read_csv(filename)
print df
```

Why store it this way?

- Different type
- Different metric

1.3.3 Converting the DateTimeUTC Column

NumPy datetime64 dtype

1.3.4 Converting the Xcoord and Ycoord

thousands=','

```
In [4]: ?pd.read_csv
```

```
In [6]: filename = os.path.join('Walrus_Data', 'Walruses.csv')
        df = pd.read_csv(filename, parse_dates=[1],
                          thousands=',')
        df.head(5)
```

```
Out[6]:
```

	Walrus	DateTimeUTC	Xcoord	Ycoord	Behav	Longitude	\
0	271	2008-05-31 19:25:00	95616.95	-528324.60	1.00900	-167.956095	
1	271	2008-06-01 03:24:00	84741.71	-511653.75	1.00050	-168.177987	
2	271	2008-06-01 11:24:00	71834.45	-491176.95	1.00625	-168.444360	
3	271	2008-06-01 19:24:00	65275.80	-478935.62	1.02025	-168.580284	
4	271	2008-06-02 03:24:00	69343.24	-473948.91	1.00775	-168.489215	

```

    Latitude
0  65.248715
1  65.401217
2  65.587969
3  65.699143
4  65.742984

```

```
In [10]: df.describe()
```

```

Out[10]:
    Walrus      Xcoord      Ycoord      Behav      Longitude \
count  454.000000    454.000000    454.000000    454.000000    454.000000
mean   281.759912   192488.612731   26570.073264    1.432807   -153.571969
std     19.599907   192785.195478   225551.527159    0.391243    59.339289
min     271.000000  -473786.460000  -528324.600000    1.000000   -179.603905
25%     271.000000    73500.247500   -74775.707500    1.056688   -167.684144
50%     271.000000   269021.360000   102471.775000    1.296250   -162.206341
75%     281.000000   291732.090000   212473.185000    1.866563   -161.616675
max     322.000000   504837.070000   246443.380000    1.996250   179.723232

```

```

    Latitude
count  454.000000
mean    70.085316
std     1.933906
min     65.248715
25%     69.287408
50%     70.793950
75%     71.576913
max     72.041307

```

1.3.5 Indexing

```
In [45]: df[2:10]
```

```

Out[45]:
    Walrus      DateTimeUTC      Xcoord      Ycoord      Behav      Longitude \
2    271  2008-06-01 11:24:00   71834.45  -491176.95    1.00625   -168.444360
3    271  2008-06-01 19:24:00   65275.80  -478935.62    1.02025   -168.580284
4    271  2008-06-02 03:24:00   69343.24  -473948.91    1.00775   -168.489215
5    271  2008-06-02 11:24:00   72634.53  -457308.67    1.00000   -168.408244
6    271  2008-06-02 19:24:00   73253.86  -425586.14    1.00000   -168.376483
7    271  2008-06-03 03:24:00   79223.97  -401784.87    1.00975   -168.229162
8    271  2008-06-03 11:24:00   77052.23  -382920.49    1.20225   -168.265869
9    271  2008-06-03 19:24:00   73380.11  -379615.25    1.24225   -168.346486

```

```

    Latitude
2  65.587969
3  65.699143
4  65.742984
5  65.891423
6  66.175663
7  66.387542
8  66.557191
9  66.587727

```

```
In [7]: print(len(df))
df[-5:]
```

454

```
Out[7]:
```

	Walrus	DateTimeUTC	Xcoord	Ycoord	Behav	Longitude
449	322	2009-07-05 12:11:00	-465471.95	169270.91	1.32675	177.051563
450	322	2009-07-05 20:11:00	-473786.46	174848.10	1.52075	176.794183
451	322	2009-07-06 04:11:00	-462401.15	175580.80	1.47450	177.097679
452	322	2009-07-06 12:11:00	-449812.23	178045.34	1.47200	177.424821
453	322	2009-07-06 20:11:00	-443963.94	180789.20	1.46775	177.568265

	Latitude
449	71.071287
450	71.104176
451	71.132315
452	71.177357
453	71.212108

1.3.6 Hierarchical columns

```
In [8]: wd = df.pivot(index='DateTimeUTC', columns='Walrus') #row, column, values (optional)
wd[:7]
```

```
Out[8]:
```

		Xcoord		Ycoord	
Walrus		271	281 322	271	281 322
DateTimeUTC					
2008-05-31 19:25:00		95616.95	NaN NaN	-528324.60	NaN NaN
2008-06-01 03:24:00		84741.71	NaN NaN	-511653.75	NaN NaN
2008-06-01 11:24:00		71834.45	NaN NaN	-491176.95	NaN NaN
2008-06-01 19:24:00		65275.80	NaN NaN	-478935.62	NaN NaN
2008-06-02 02:25:00		NaN	65600.95 NaN	NaN	-417464.74 NaN
2008-06-02 03:24:00		69343.24	NaN NaN	-473948.91	NaN NaN
2008-06-02 10:24:00		NaN	61574.27 NaN	NaN	-421676.30 NaN

		Behav		Longitude	
Walrus		271	281 322	271	281 322
DateTimeUTC					
2008-05-31 19:25:00		1.00900	NaN NaN	-167.956095	NaN NaN
2008-06-01 03:24:00		1.00050	NaN NaN	-168.177987	NaN NaN
2008-06-01 11:24:00		1.00625	NaN NaN	-168.444360	NaN NaN
2008-06-01 19:24:00		1.02025	NaN NaN	-168.580284	NaN NaN
2008-06-02 02:25:00		NaN	1.544 NaN	NaN	-168.541792 NaN
2008-06-02 03:24:00		1.00775	NaN NaN	-168.489215	NaN NaN
2008-06-02 10:24:00		NaN	1.567 NaN	NaN	-168.633333 NaN

		Latitude	
Walrus		271	281 322
DateTimeUTC			
2008-05-31 19:25:00		65.248715	NaN NaN
2008-06-01 03:24:00		65.401217	NaN NaN
2008-06-01 11:24:00		65.587969	NaN NaN
2008-06-01 19:24:00		65.699143	NaN NaN
2008-06-02 02:25:00		NaN	66.250190 NaN
2008-06-02 03:24:00		65.742984	NaN NaN
2008-06-02 10:24:00		NaN	66.213262 NaN

```
In [48]: wd['Behav'][:5]
```


[illegible]

[illegible]

1.3.7 Extracting with a Condition

- Extracting Walrus 271 from the table

```
In [18]: df[df.Walrus == 271][:5]
```

```
Out[18]:
```

	Walrus	DateTimeUTC	Xcoord	Ycoord	Behav	Longitude
0	271	2008-05-31 19:25:00	95616.95	-528324.60	1.00900	-167.956095
1	271	2008-06-01 03:24:00	84741.71	-511653.75	1.00050	-168.177987
2	271	2008-06-01 11:24:00	71834.45	-491176.95	1.00625	-168.444360
3	271	2008-06-01 19:24:00	65275.80	-478935.62	1.02025	-168.580284
4	271	2008-06-02 03:24:00	69343.24	-473948.91	1.00775	-168.489215

	Latitude
0	65.248715
1	65.401217
2	65.587969
3	65.699143
4	65.742984

```
In [21]: wd.columns
```

```
Out[21]: MultiIndex(levels=[[u'Xcoord', u'Ycoord', u'Behav', u'Longitude', u'Latitude'], [271, 281, 322],
labels=[[0, 0, 0, 1, 1, 1, 2, 2, 2, 3, 3, 3, 4, 4, 4], [0, 1, 2, 0, 1, 2, 0, 1, 2, 0, 1, 2, 0, 1, 2], [0, 1, 2, 0, 1, 2, 0, 1, 2, 0, 1, 2, 0, 1, 2]],
names=[None, u'Walrus'])
```

```
In [32]: wd[(['Latitude', 271), ('Longitude', 271)][:5]
```

```
Out[32]:
```


		Latitude	Longitude
Walrus		271	271
DateTimeUTC			
2008-05-31 19:25:00		65.248715	-167.956095
2008-06-01 03:24:00		65.401217	-168.177987
2008-06-01 11:24:00		65.587969	-168.444360
2008-06-01 19:24:00		65.699143	-168.580284
2008-06-02 02:25:00		NaN	NaN

1.4 Simple Plotting


- Plot the behavior of each walrus over time

```
In [34]: df.plot()
```

```
Out[34]: <matplotlib.axes._subplots.AxesSubplot at 0x1057dcd50>
```



```
In [44]: wd['Behav'].plot(style=['bo','rx','gv'])
plt.ylabel('behavior')
plt.title('Walrus plotting exercise')
plt.savefig('walrus_behav.pdf')
plt.show()
```

In [61]: !ls

```
01_exploratory_pandas.ipynb      06_CSV_to_NetCDF_Solution.ipynb
01_introduction-IPython-notebook.ipynb  Walrus_Data
01_introduction-IPython-notebook.pdf    data_overview.png
01_introduction-IPython-notebook_files  ipython-notebook-keyboard.png
02_Data_Transfer.pdf                ipython-notebook-sharing.png
03_HPC_File_Systems.pdf              ipython-notebook.png
04_python_reading-plotting.ipynb      rc_logo.png
05_Data_Conversion_Cleaning.pdf        traditional_python.png
06_CSV_to_NetCDF_Exercise.ipynb        walrus_behav.png
```

1.4.1 Importing the image into the Markdown

In []: